

DS502- HW4

Mahdi Alouane and Rahul Pande

1. (10 points) Section 6.8, page 259, question 2

- (a) iii. holds true. Lasso puts a budget constraint on the parameters which decreases the model variance and it reduces overfitting. However, when we put a constraint, the model bias increases. From the bias-variance trade-off concept we can say that the lasso regression will give better prediction when its increase in bias is less than its decrease in variance.
- (b) iii. holds true. Like Lasso above, Ridge also puts a budget constraint on the parameters which decreases the model variance and it reduces overfitting. However, when we put a constraint, the model bias increases. Similarly as above, from the bias-variance trade-off concept we can say that the Ridge regression will give better prediction when its increase in bias is less than its decrease in variance.
- (c) ii. holds true. Since non-linear methods are more flexible, they have higher variance than least squares regression but lower bias. Again, from the bias variance trade-off, if the increase on model variance is less than the decrease in bias, then non-linear model will have better prediction accuracy.

2. (20 points) Section 6.8, page 264, question 11

(a)

```
library(MASS)
library(leaps)
library(glmnet)
```

```
## Loading required package: Matrix
## Loading required package: foreach
## Loaded glmnet 2.0-16
```

```
attach(Boston)
```

```
colSums(sapply(Boston, is.na))
```

```
##      crim      zn      indus      chas      nox      rm      age      dis      rad
##         0         0         0         0         0         0         0         0         0
##      tax ptratio  black  lstat  medv
##         0         0         0         0         0
```

```
# No NAs in the dataset
```

```
# k-fold cross validation
```

```
k = 10
```

```
n = dim(Boston)[1]
```

```
p = dim(Boston)[2]-1
```

```
set.seed(123)
```

```
folds = sample(rep(1:k, length = nrow(Boston)), replace = T)
```

```

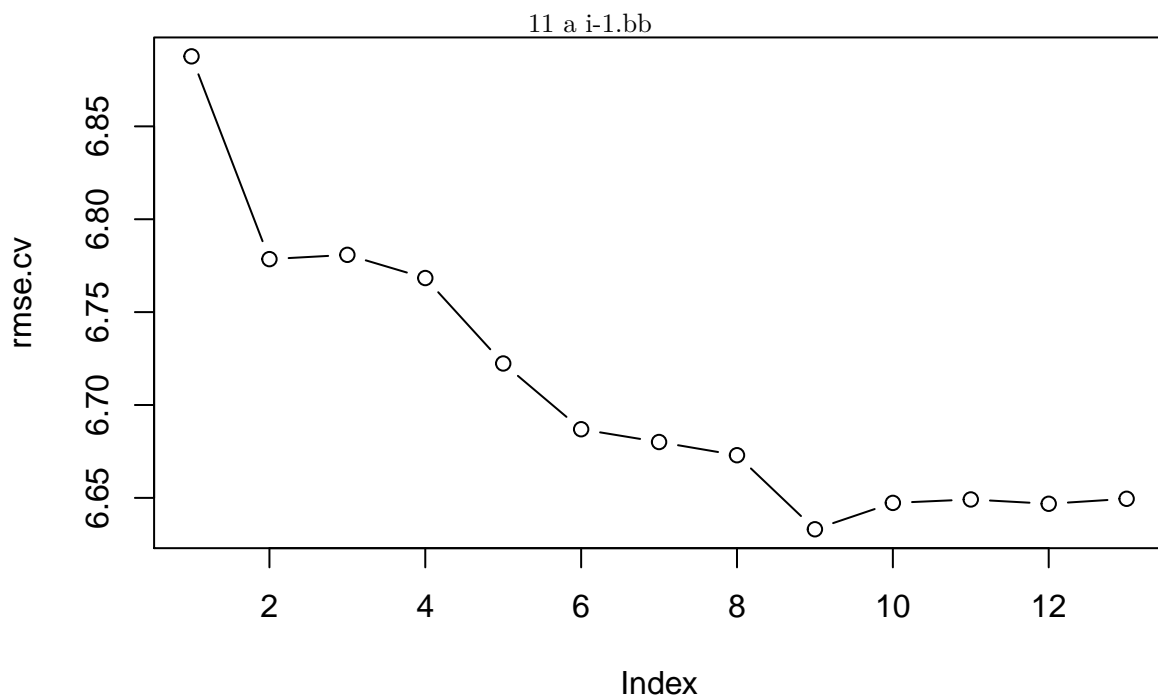
form.subset = as.formula("crim ~ .")

cv.errors = matrix(NA, k, p)

for (i in 1:k) {
  subset.fit <- regsubsets(form.subset, Boston[folds!=i, ], nvmax = p)
  for (n.subset in 1:p) {
    m.mat <- model.matrix(form.subset, Boston[folds==i, ])
    best.coef <- coef(subset.fit, id=n.subset)
    pred <- m.mat[, names(best.coef)] %*% best.coef
    # mean squared error
    error = mean((Boston[folds==i, ]$crim - pred)^2)
    cv.errors[i,n.subset] = error
  }
}

# root mean squared values
rmse.cv = sqrt(apply(cv.errors, 2, mean))
plot(rmse.cv, type = "b")

```



```

which(rmse.cv == min(rmse.cv))

## [1] 9

rmse.cv[which(rmse.cv == min(rmse.cv))]

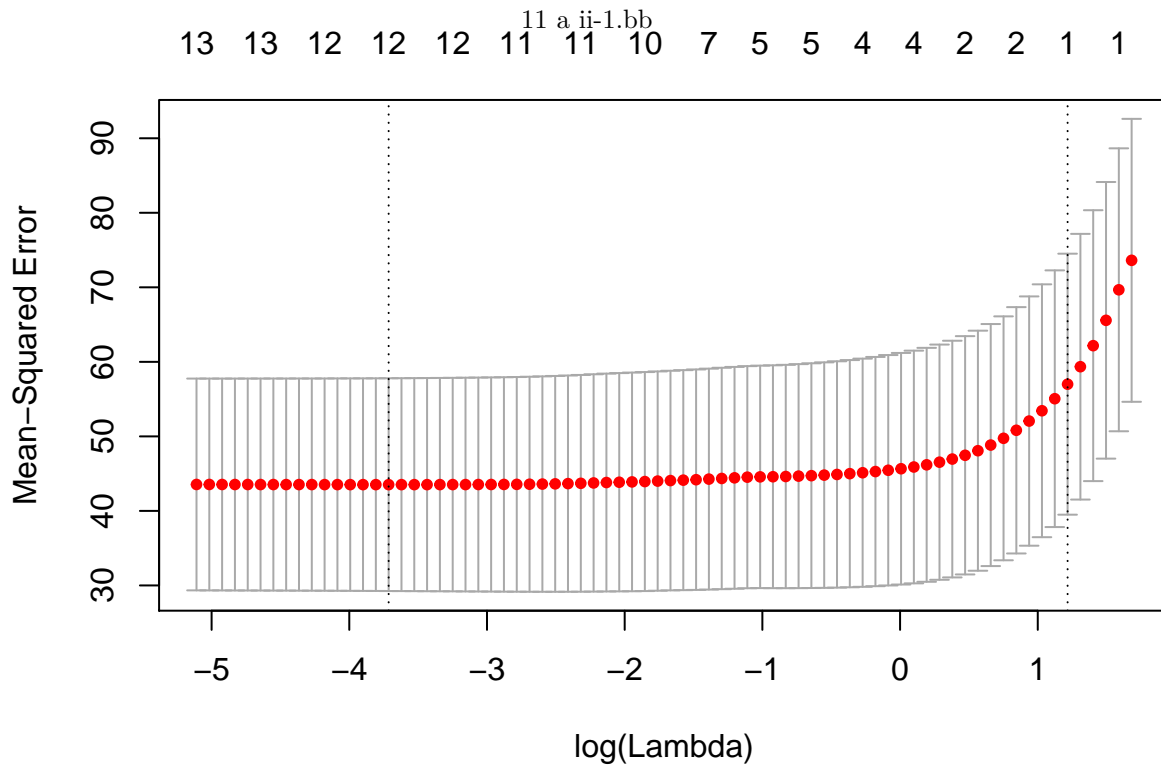
## [1] 6.633116

# Lasso regression

X = model.matrix(crim ~ ., data = Boston)[, -1]
y = Boston$crim
cv.lasso = cv.glmnet(X, y, type.measure = "mse", nfolds = 10)

```

```
plot(cv.lasso)
```



```
coef(cv.lasso)[,1]
```

```
## (Intercept)      zn      indus      chas      nox      rm
##  1.4186415  0.0000000  0.0000000  0.0000000  0.0000000  0.0000000
##      age      dis      rad      tax      ptratio  black
##  0.0000000  0.0000000  0.2298449  0.0000000  0.0000000  0.0000000
##      lstat      medv
##  0.0000000  0.0000000
```

```
# One standard error lamda to avoid overfitting
```

```
chosen.lambda = cv.lasso$lambda.1se
```

```
# root mean square error for the chosen lamdba
```

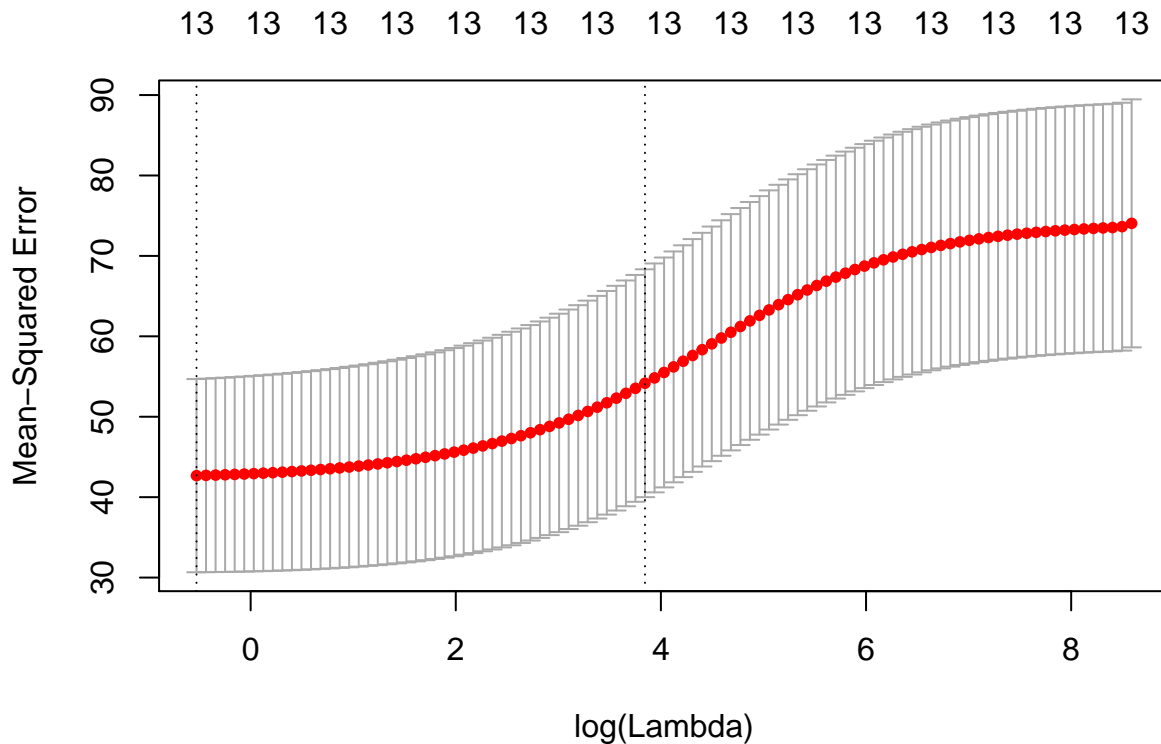
```
sqrt(cv.lasso$cvm[cv.lasso$lambda == chosen.lambda])
```

```
## [1] 7.549995
```

```
cv.ridge = cv.glmnet(X, y, type.measure = "mse", alpha = 0, nfolds = 10)
```

```
plot(cv.ridge)
```

11 a iii-1.bb



```
coef(cv.ridge)
```

```
## 14 x 1 sparse Matrix of class "dgCMatrix"
##              1
## (Intercept)  0.757566248
## zn          -0.002519179
## indus       0.036530724
## chas        -0.259774554
## nox         2.423435801
## rm          -0.169106780
## age         0.007845252
## dis         -0.125063596
## rad         0.067739853
## tax         0.002972047
## ptratio     0.093676787
## black       -0.003748456
## lstat       0.049092196
## medv        -0.032058198
```

```
# One standard error lamda to avoid overfitting
```

```
chosen.lambda = cv.ridge$lambda.1se
```

```
# root mean square error for the chosen lamdba
```

```
sqrt(cv.ridge$cvm[cv.ridge$lambda == chosen.lambda])
```

```
## [1] 7.359946
```

```
library(pls)
```

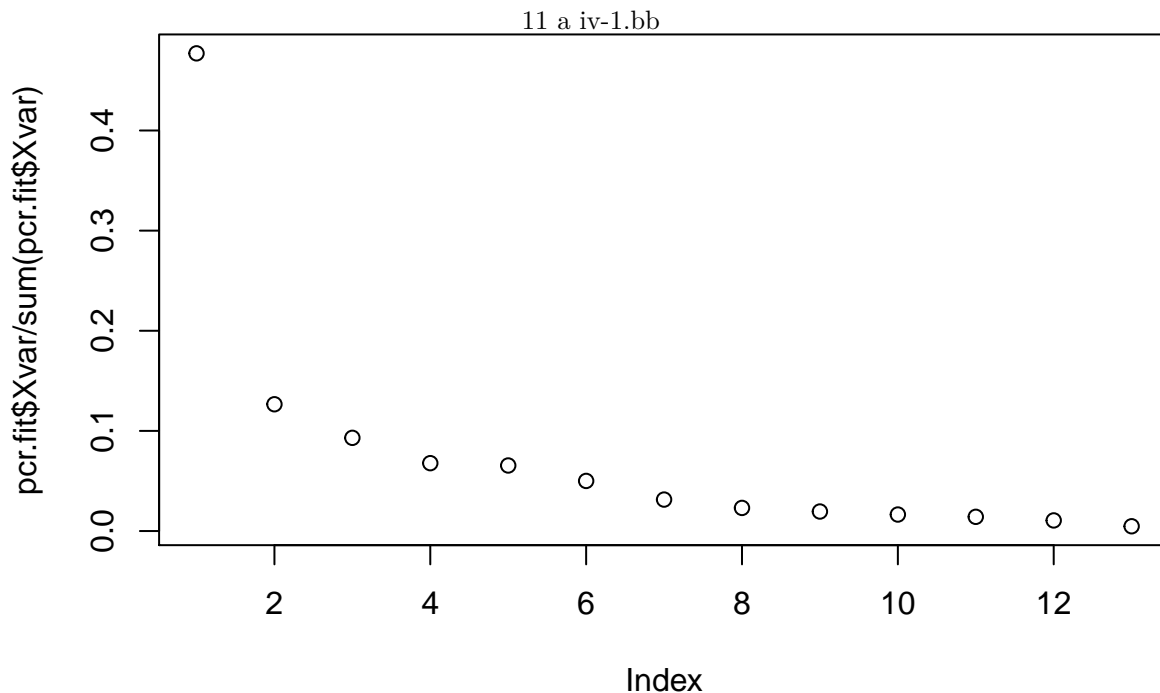
```
##
```

```
## Attaching package: 'pls'
```

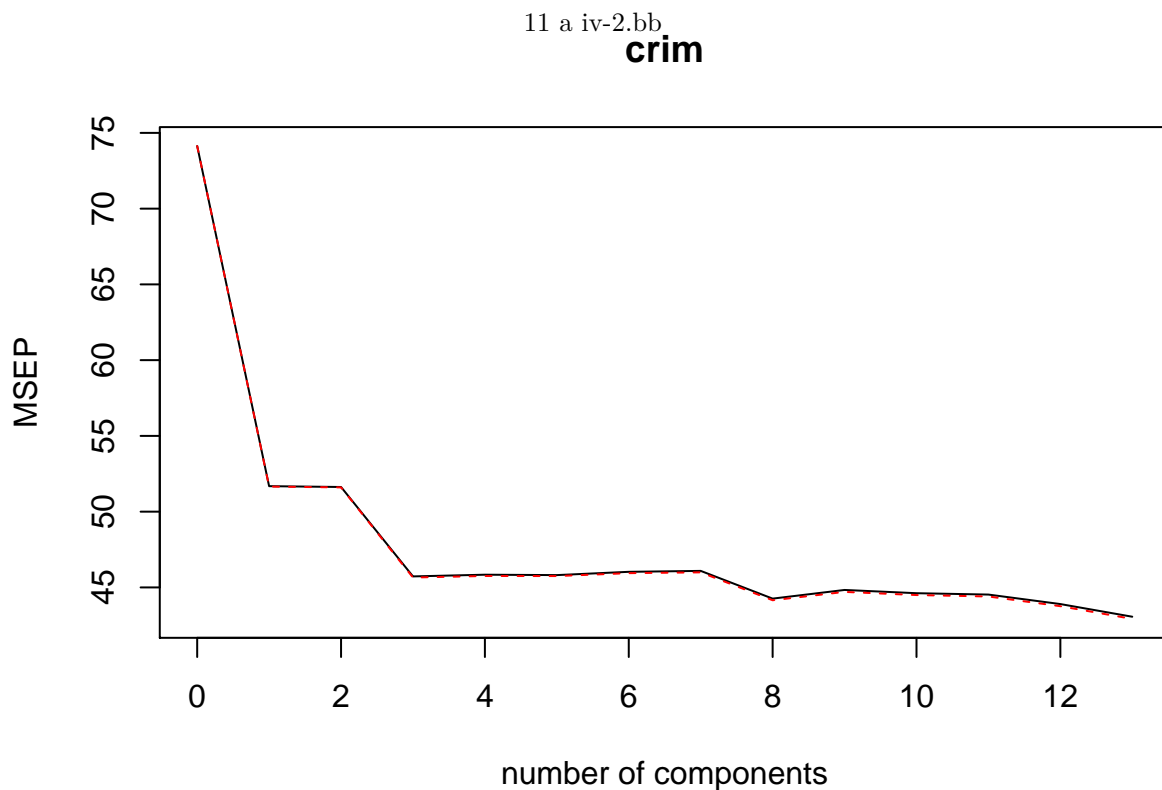
```
## The following object is masked from 'package:stats':
##
##   loadings
pcr.fit = pcr(crim ~ ., data = Boston, scale = TRUE, validation = "CV", segments= 10)
summary(pcr.fit)

## Data:      X dimension: 506 13
## Y dimension: 506 1
## Fit method: svdpc
## Number of components considered: 13
##
## VALIDATION: RMSEP
## Cross-validated using 10 random segments.
##      (Intercept)  1 comps  2 comps  3 comps  4 comps  5 comps  6 comps
## CV              8.61    7.189    7.185    6.762    6.770    6.769    6.784
## adjCV           8.61    7.187    7.183    6.757    6.764    6.764    6.778
##      7 comps  8 comps  9 comps 10 comps 11 comps 12 comps 13 comps
## CV          6.789    6.653    6.696    6.680    6.673    6.626    6.563
## adjCV        6.782    6.645    6.686    6.671    6.663    6.615    6.551
##
## TRAINING: % variance explained
##      1 comps  2 comps  3 comps  4 comps  5 comps  6 comps  7 comps
## X          47.70    60.36    69.67    76.45    82.99    88.00    91.14
## crim       30.69    30.87    39.27    39.61    39.61    39.86    40.14
##      8 comps  9 comps 10 comps 11 comps 12 comps 13 comps
## X          93.45    95.40    97.04    98.46    99.52    100.0
## crim       42.47    42.55    42.78    43.04    44.13    45.4

# variance explanation
plot(pcr.fit$Xvar/sum(pcr.fit$Xvar))
```



```
# validation plot
validationplot(pcr.fit, val.type='MSEP')
```



```
ncomp = 4

set.seed(123)
folds = sample(rep(1:k, length = nrow(Boston)), replace = T)

cv.errors = rep(0, k)

for (i in 1:k) {
  pcr.fit <- pcr(crim ~ ., data = Boston[folds != i, ], scale = TRUE)
  pred = predict(pcr.fit, Boston[folds == i, ], ncomp=ncomp)
  error = mean((Boston[folds == i, ]$crim - pred)^2)
  cv.errors[i] = error
}
sqrt(mean(cv.errors))
```

```
## [1] 6.875975
```

(b)

```
results <- rbind(
  c("Best Subset", 6.633116, 9),
  c("Lasso Regression", 7.549995, 1),
  c("Ridge Regression", 7.359946, 13),
  c("PCR", 6.875975, 4)
)
colnames(results) <- c("Method", "MSE", "# predictors")
```

```
knitr::kable(results)
```

Method	MSE	# predictors
Best Subset	6.633116	9
Lasso Regression	7.549995	1
Ridge Regression	7.359946	13
PCR	6.875975	4

From the above table, we chose PCR model with 4 predictors as it has the mse very close to the lowest mse and is simpler model than the best subset model, since it has 4 predictors against the 9 predictors and thus has lower chances of overfitting the data. Second choice would be Best Subset model with 9 predictors as it has the lowest cross validation mse and has less number of predictors than Ridge Regression. Lasso Regression here seems to be underfit in this case.

- (c) No. PCR has only 4 predictors since from the graph, after 4 predictors, adding another predictor does not increase the explained variance a lot. Hence we have taken only 4 predictors. We can also have criterion like taking n components which explain at least $x\%$ of the variance.

3. (10 points) Section 7.9, Page 298, question 3

We have $b_1(X) = X, b_2(X) = (X - 1)^2 * I(X \geq 1)$

For $X \geq 1, I(X \geq 1) = 1$, and $X < 1, I(X \geq 1) = 0$

Substituting $\hat{\beta}_0 = 1, \hat{\beta}_1 = 1, \hat{\beta}_2 = 2$ in

$$Y = \beta_0 + \beta_1 b_1(X) + \beta_2 b_2(X) + \epsilon$$

We get,

$$\hat{Y} = 1 + b_1(X) - 2b_2(X)$$

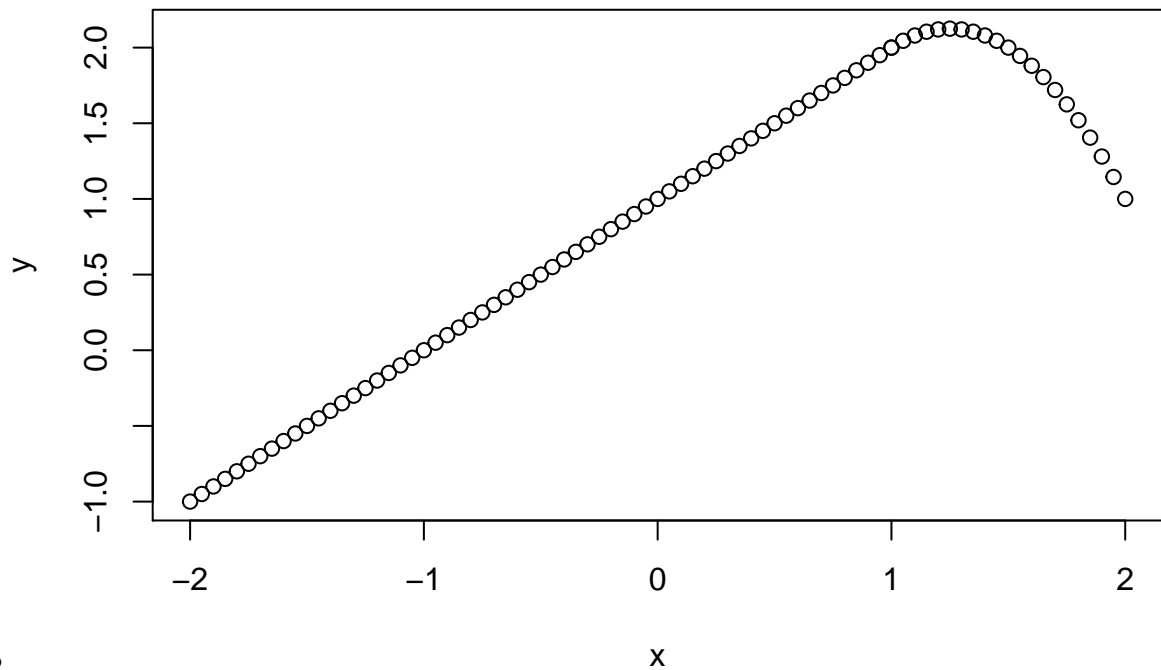
For $X \geq 1, \hat{Y} = 1 + X - 2(X - 1)^2 = -1 + 5X - 2X^2$, and for $X < 1, \hat{Y} = 1 + X$

```
x_lower = seq(-2, 1, by = 0.05)
x_upper = seq(1, 2, by = 0.05)

y_lower = 1 + x_lower
y_upper = -1 + 5 * x_upper - 2 * (x_upper ^ 2)

x <- c(x_lower, x_upper)
y <- c(y_lower, y_upper)

plot(x,y)
```



2-1.bb

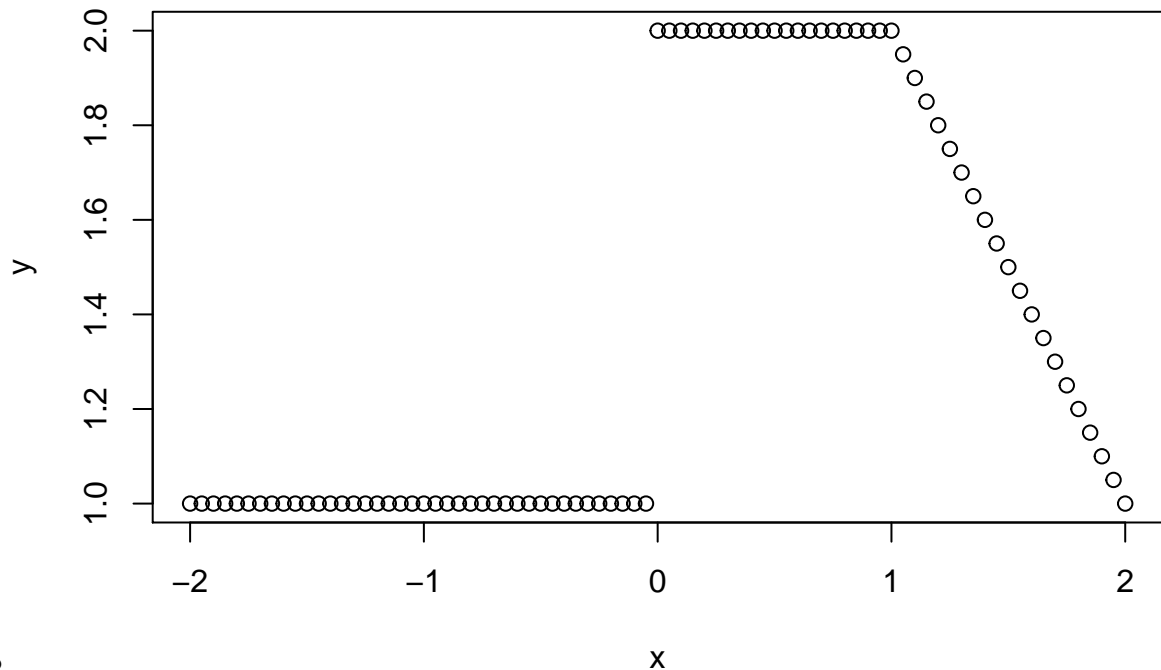
$Y = 1 + X$, for $X < 1$ Now, for $X = 0$, $Y = 1$

Therefore y-intercept is 1. Slope is 1 when $X < 1$ and by taking derivative for Y where $X \geq 1$, we get slope as $5 - 4X$

4. (10 points) Section 7.9, Page 298, question 4

Similar from above, we can split the function into multiple domains. Since there are a lot of cuts in this, we use the I function in R to enforce the conditions on x. It is as below.

```
x = seq(-2, 2, 0.05)
y = 1 + 1 * I(x <= 2 & x >= 0) - (x-1) * I(x <= 2 & x >= 1) + 3 * (x-3) * I(x <= 4 & x >= 3) + I(x <=
plot(x, y)
```

```
4-1.bb
# y-intercept
y[which(x==0)]
```

```
## [1] 2
```

The y-intercept is 2 (`y[which(x==0)]`). Slope is -1 for $1 \leq X \leq 2$ and 0 for $-2 \leq X < 0$ and $0 < X \leq 1$. The function is discontinuous at $x=0$

5. (10 points) Section 7.9, Page 299, question 6

a.

First, we start by performing a 10-fold cross validation as shown below:

```
# Import the required libraries
library(ISLR)
library(boot)

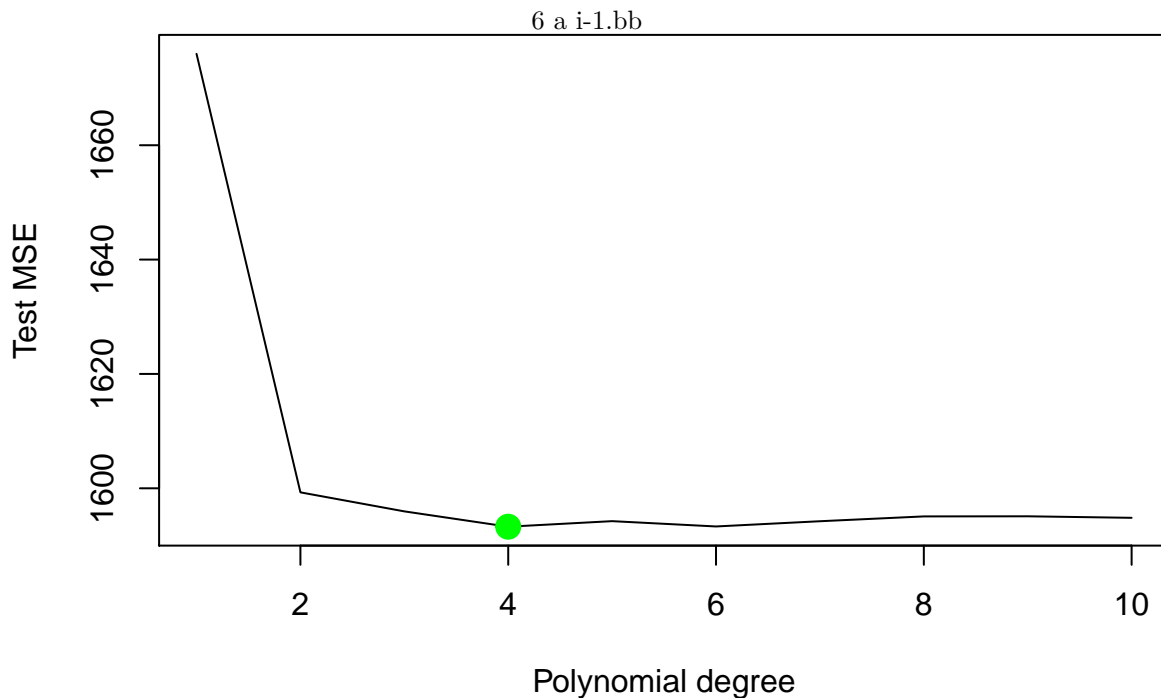
# Fix the random seed
set.seed(4)

# Initialize the error/degree vector
errors = rep(NA, 10)

# For each degree from 1 to 10
for (d in 1:10) {
  # Fit a polynomial model of degree d
  fit = glm(wage ~ poly(age, d), data = Wage)
  # Estimate of the test MSE with 10-fold cross validation
  errors[d] = cv.glm(Wage, fit, K = 10)$delta[1]
}

# Plot MSE vs. Degree
plot(1:10, errors, xlab = "Polynomial degree", ylab = "Test MSE", type = "l")
```

```
# Highlight the lowest test MSE value in the plot
points(which.min(errors), errors[which.min(errors)], col = "green", cex = 2.5, pch = 20)
```



We can see that the polynomial model with $d = 4$ outperforms the other models on the testing set which corresponds to the optimal model. In the section below, we use ANOVA for testing the null hypothesis that this model is complex enough to explain the data vs. the alternative hypothesis that a more complex model is required.

```
# We fit different polynomial models going from degree 1 to 10
fit1 = lm(wage ~ age, data = Wage)
fit2 = lm(wage ~ poly(age, 2), data = Wage)
fit3 = lm(wage ~ poly(age, 3), data = Wage)
fit4 = lm(wage ~ poly(age, 4), data = Wage)
fit5 = lm(wage ~ poly(age, 5), data = Wage)
fit6 = lm(wage ~ poly(age, 6), data = Wage)
fit7 = lm(wage ~ poly(age, 7), data = Wage)
fit8 = lm(wage ~ poly(age, 8), data = Wage)
fit9 = lm(wage ~ poly(age, 9), data = Wage)
fit10 = lm(wage ~ poly(age, 10), data = Wage)

# We use the null hypothesis test ANOVA
anova(fit1, fit2, fit3, fit4, fit5, fit6, fit7, fit8, fit9, fit10)
```

```
## Analysis of Variance Table
##
## Model 1: wage ~ age
## Model 2: wage ~ poly(age, 2)
## Model 3: wage ~ poly(age, 3)
## Model 4: wage ~ poly(age, 4)
## Model 5: wage ~ poly(age, 5)
## Model 6: wage ~ poly(age, 6)
## Model 7: wage ~ poly(age, 7)
```

```
## Model 8: wage ~ poly(age, 8)
## Model 9: wage ~ poly(age, 9)
## Model 10: wage ~ poly(age, 10)
##      Res.Df      RSS Df Sum of Sq      F      Pr(>F)
## 1      2998 5022216
## 2      2997 4793430  1    228786 143.7638 < 2.2e-16 ***
## 3      2996 4777674  1     15756  9.9005  0.001669 **
## 4      2995 4771604  1      6070  3.8143  0.050909 .
## 5      2994 4770322  1      1283  0.8059  0.369398
## 6      2993 4766389  1      3932  2.4709  0.116074
## 7      2992 4763834  1      2555  1.6057  0.205199
## 8      2991 4763707  1       127  0.0796  0.777865
## 9      2990 4756703  1      7004  4.4014  0.035994 *
## 10     2989 4756701  1         3  0.0017  0.967529
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

We can observe that a polynomial fit with degree 4 provides a statistically significant fit to the data (degree 3 too), however, the other models are not justified since they provide higher p-values. We can say the the ANOVA test proves the results found in the plot above.

In this section, we fit the data with the previous model:

```
# Plot the Wage vs. Age data
plot(wage ~ age, data = Wage)

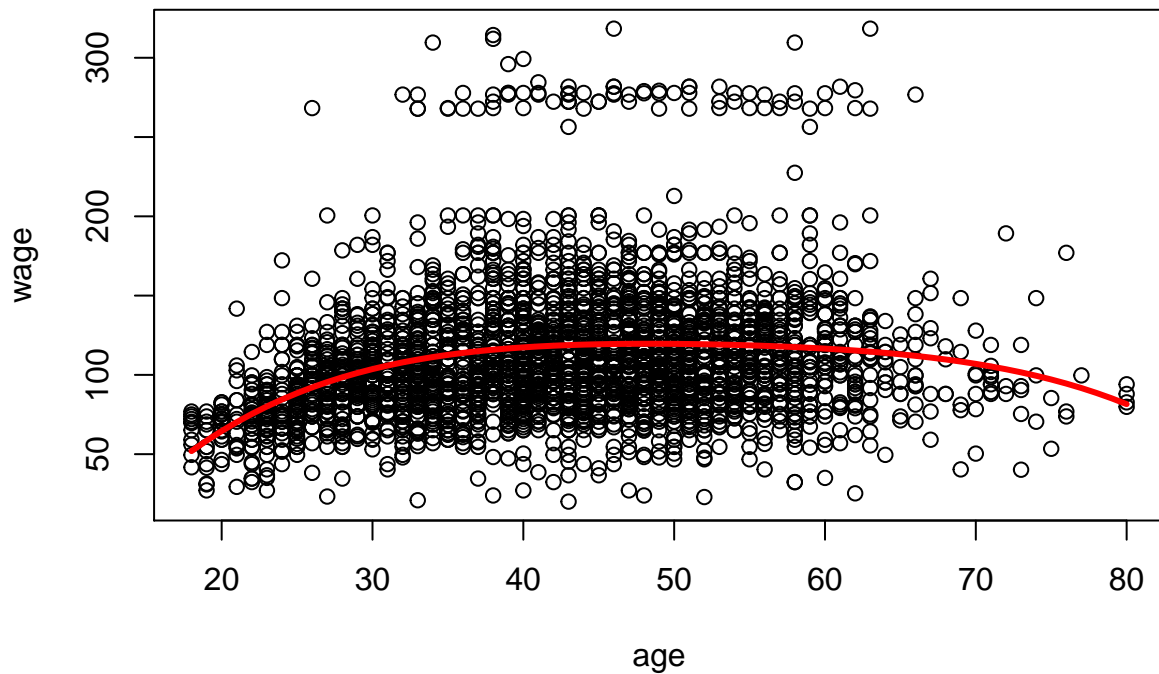
# Calculate the age grid (lower and upper limits)
limits = range(Wage$age)
ageGrid = seq(from = limits[1], to = limits[2])

# Fit the data with the quartic polynomial model
fit = lm(wage ~ poly(age, 4), data = Wage)

# Calculate the predictions for this model
predictions = predict(fit, newdata = list(age = ageGrid))

# Display the model on the plot
lines(ageGrid, predictions, col = "red", lwd = 3)
```

6 a iii-1.bb

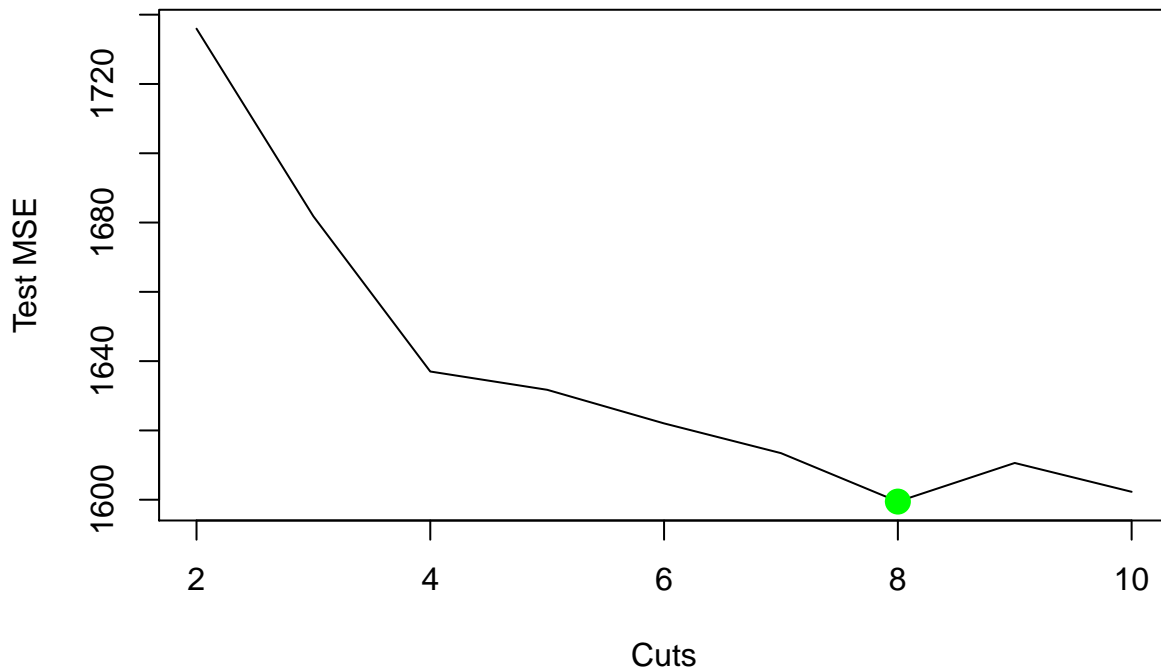


b.

We repeat the same process for the step function using 10-fold cross validation.

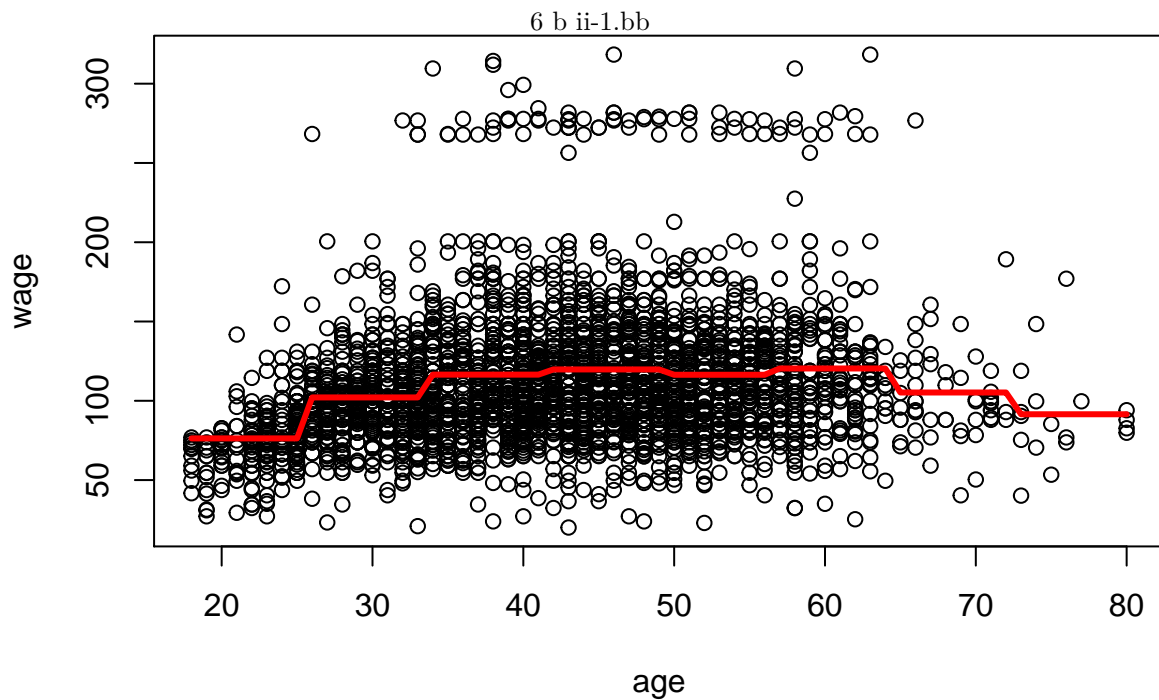
```
errors <- rep(NA, 10)
for (intervals in 2:10) {
  Wage$age.cut = cut(Wage$age, intervals)
  fit = glm(wage ~ age.cut, data = Wage)
  errors[intervals] = cv.glm(Wage, fit, K = 10)$delta[1]
}
plot(2:10, errors[-1], xlab = "Cuts", ylab = "Test MSE", type = "l")
points(which.min(errors), errors[which.min(errors)], col = "green", cex = 2.5, pch = 20)
```

6 b i-1.bb



We can observe that the optimal number of `cuts` = 8 since it corresponds to the lowest error on the testing set. In the section below, we plot the data and the 8-cuts step model fit.

```
plot(wage ~ age, data = Wage)
agelims = range(Wage$age)
age.grid = seq(from = agelims[1], to = agelims[2])
fit = glm(wage ~ cut(age, 8), data = Wage)
preds = predict(fit, data.frame(age = age.grid))
lines(age.grid, preds, col = "red", lwd = 3)
```



6. (20 points) Section 7.9, Page 299, question 7

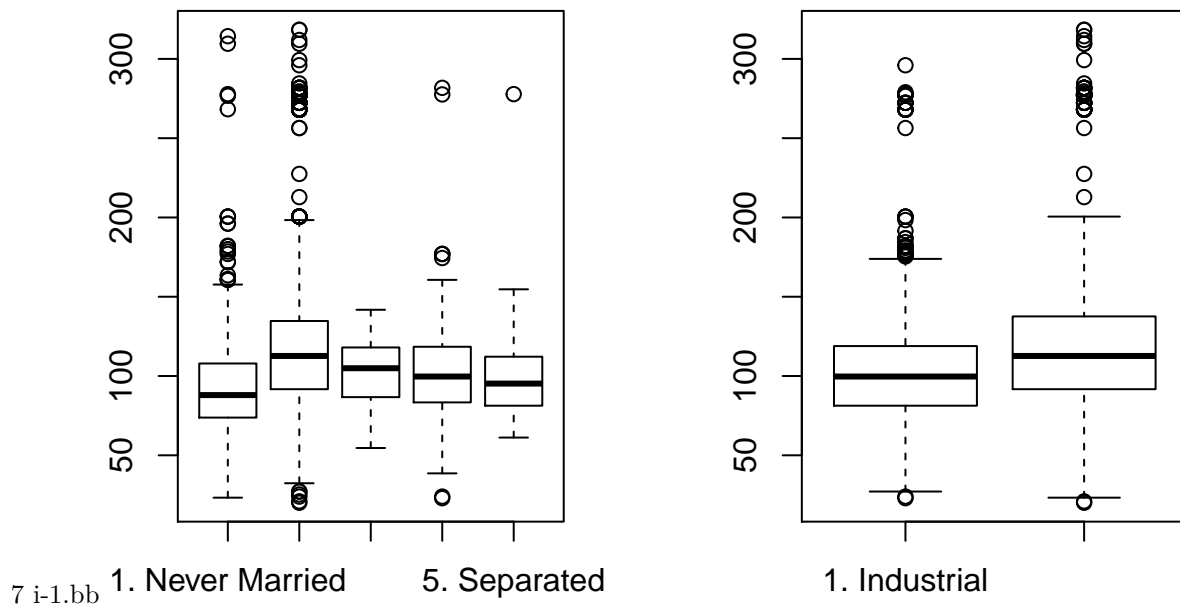
Let's first begin by exploring the relationship between `wage` from a side and `maritl` and `jobclass` from the other side.

```
# Fix the random seed
set.seed(1)

# Display the summary of the features maritl and jobclass (index of values)
summary(Wage[, c("maritl", "jobclass")])
```

```
##                maritl                jobclass
## 1. Never Married: 648    1. Industrial :1544
## 2. Married       :2074    2. Information:1456
## 3. Widowed       :   19
## 4. Divorced      :  204
## 5. Separated     :   55
```

```
# Display the wage vs. maritl and wage vs. jobclass side by side
par(mfrow = c(1, 2))
plot(Wage$maritl, Wage$wage)
plot(Wage$jobclass, Wage$wage)
```



From the plots above, we can see that a `married` marital status corresponds to higher wages (followed by widowed) and job class corresponding to `informational` presents higher wages too on an average.

In this section, let's experiment different models using natural spline functions of the variables `maritl`, `jobclass`, `year`, `education` and `age`.

```
library(gam)
```

```
## Loading required package: splines
```

```
## Loaded gam 1.16
```

```
fit0 = gam(wage ~ lo(age,span=200,degree=1), data = Wage)
fit1 = gam(wage ~ lo(age,span=200,degree=1) + year, data = Wage)
fit2 = gam(wage ~ s(year,4)+lo(age,span=200,degree=1), data = Wage)
```

```

fit4 = gam(wage ~ s(year,4)+lo(age,span=200,degree=1) + education, data = Wage)
fit5 = gam(wage ~ s(year,4)+lo(age,span=200,degree=1) + education + jobclass, data = Wage)
fit6 = gam(wage ~ s(year,4)+lo(age,span=200,degree=1) + education + maritl, data = Wage)
fit7 = gam(wage ~ s(year,4)+lo(age,span=200,degree=1) + education + jobclass + maritl, data = Wage)
anova(fit0, fit1, fit2, fit4, fit5, fit6, fit7)

```

```
## Analysis of Deviance Table
```

```
##
```

```
## Model 1: wage ~ lo(age, span = 200, degree = 1)
```

```
## Model 2: wage ~ lo(age, span = 200, degree = 1) + year
```

```
## Model 3: wage ~ s(year, 4) + lo(age, span = 200, degree = 1)
```

```
## Model 4: wage ~ s(year, 4) + lo(age, span = 200, degree = 1) + education
```

```
## Model 5: wage ~ s(year, 4) + lo(age, span = 200, degree = 1) + education +
```

```
##      jobclass
```

```
## Model 6: wage ~ s(year, 4) + lo(age, span = 200, degree = 1) + education +
```

```
##      maritl
```

```
## Model 7: wage ~ s(year, 4) + lo(age, span = 200, degree = 1) + education +
```

```
##      jobclass + maritl
```

```
##      Resid. Df Resid. Dev      Df Deviance  Pr(>Chi)
```

```
## 1      2998      5022190
```

```
## 2      2997      5004580 1.0000      17610 0.0001540 ***
```

```
## 3      2994      4999892 2.9996       4689 0.2822804
```

```
## 4      2990      3849469 4.0000     1150423 < 2.2e-16 ***
```

```
## 5      2989      3835853 1.0000       13616 0.0008753 ***
```

```
## 6      2986      3685252 3.0000      150601 < 2.2e-16 ***
```

```
## 7      2985      3670120 1.0000       15132 0.0004512 ***
```

```
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

We can conclude that models 4 and 6 are statistically significant with respectively p-values equal to $< 2.2e-16$ and $< 2.2e-16$. Hence, we can say that the model gets a statistically significant improvement by including the year spline, age local regression, education and maritl. However, the model is less significant for the cases where we introduced jobclass.

In this section, we are going to plot our 7th model as shown below:

```
attach(Wage)
```

```
## The following object is masked from Boston:
```

```
##
```

```
##      age
```

```
# Calculate the prediction with our model
```

```
myPreds = predict(fit6,se=TRUE)
```

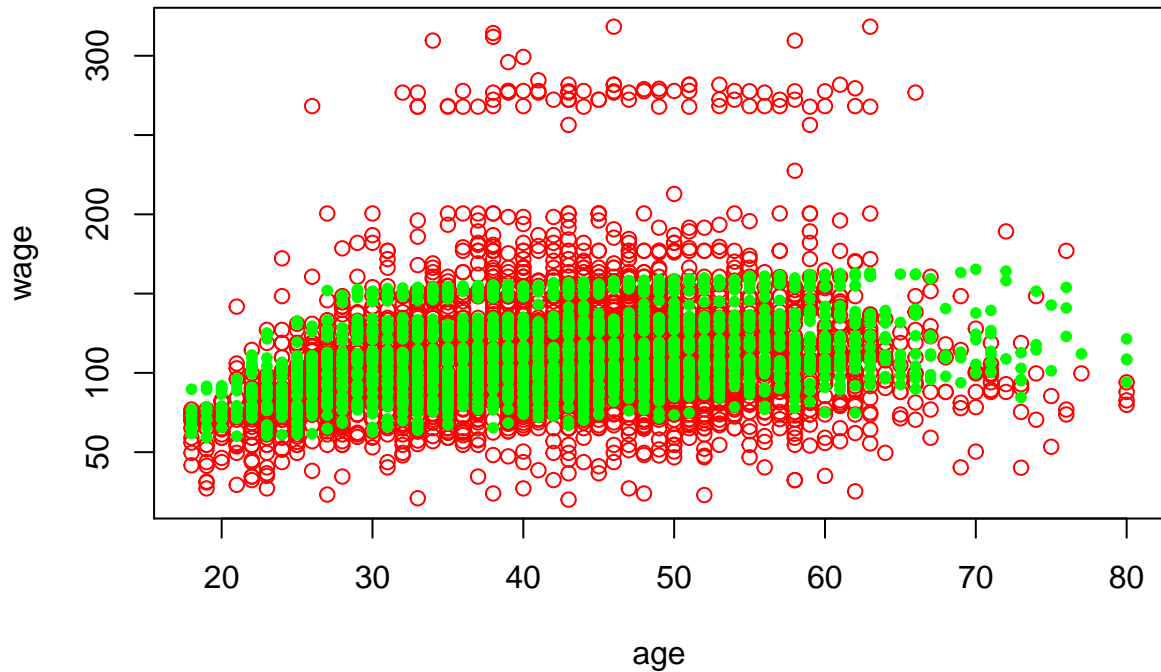
```
# Plot the data
```

```
plot(age,wage, col='red')
```

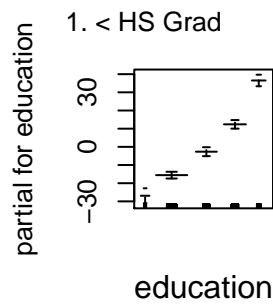
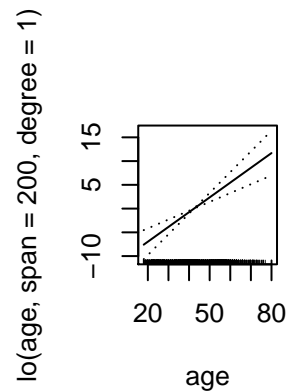
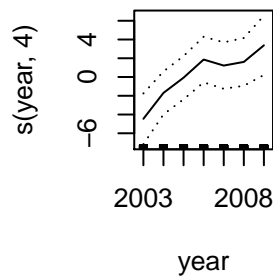
```
# Plot the predictions
```

```
points(age,myPreds$fit,col='green',pch=20)
```

7 iii-1.bb



```
# Fancy GAM plot
par(mfrow = c(2, 2), pty = "s")
plot(fit6, se=TRUE)
```



7 iii-2.bb

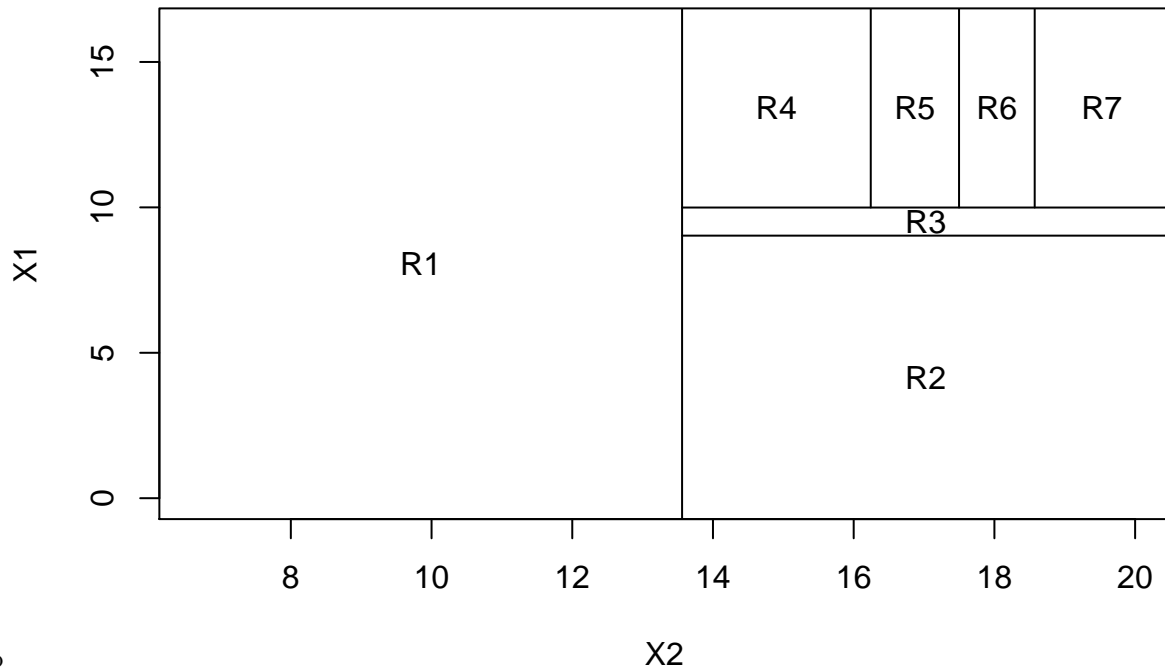
education

maritl

We can observe our data is not fitted with a linear model (plane) but instead is fitted with multiple lines each corresponding to a value of the categorical variables `education` and `maritl`. We can see that our model is better fitted by combining multiple models such as `spline`, `local regression` and `linear regression` with `continous` and `categorical` variables.

7. (10 points) Section 8.4, Page 332, question 1

```
library(tree)
syn_data <- data.frame(mvrnorm(n= 100, mu = c(10,15), matrix(c(10,3,3,5),2,2)))
color <- sample(c("red", "blue"), nrow(syn_data), replace = T)
syn_data$class <- ifelse(color == "red", 1, 0)
tree.fit <- tree(class ~ ., data = syn_data, control = tree.control(nobs = nrow(syn_data), mindev = 0.0))
partition_plot(tree.fit)
```



1-1.bb

8. (10 points) Section 8.4, Page 333-334, question 8

a.

```
library(ISLR)
set.seed(1)

# Splitting the data according to a 50/50 ratio

split = sample(1:nrow(Carseats), nrow(Carseats) / 2)
Carseats.train = Carseats[split, ]
Carseats.test = Carseats[-split, ]
```

b.

```
# Import the library
library(tree)

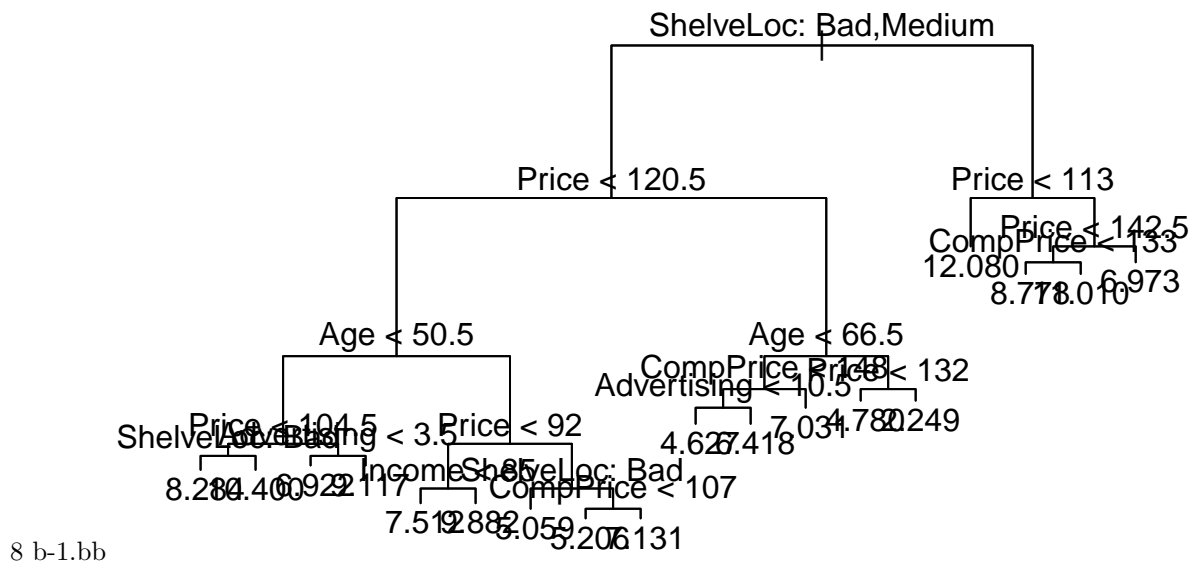
tree.carseats = tree(Sales ~ ., data = Carseats.train)
summary(tree.carseats)
```

##

Regression tree:

```
## tree(formula = Sales ~ ., data = Carseats.train)
## Variables actually used in tree construction:
## [1] "ShelveLoc" "Price" "Age" "Advertising" "Income"
## [6] "CompPrice"
## Number of terminal nodes: 18
## Residual mean deviance: 2.36 = 429.5 / 182
## Distribution of residuals:
## Min. 1st Qu. Median Mean 3rd Qu. Max.
## -4.2570 -1.0360 0.1024 0.0000 0.9301 3.9130

plot(tree.carseats)
text(tree.carseats, pretty = 0)
```



We can see that the most important variable in the tree is `ShelveLoc` and then `Price` which means that knowing the shelf location would be the information that helps us the most to predict the `sales` of the carseats.

```
pred <- predict(tree.carseats, newdata = Carseats.test)
mean((pred - Carseats.test$Sales)^2)
```

```
## [1] 4.148897
```

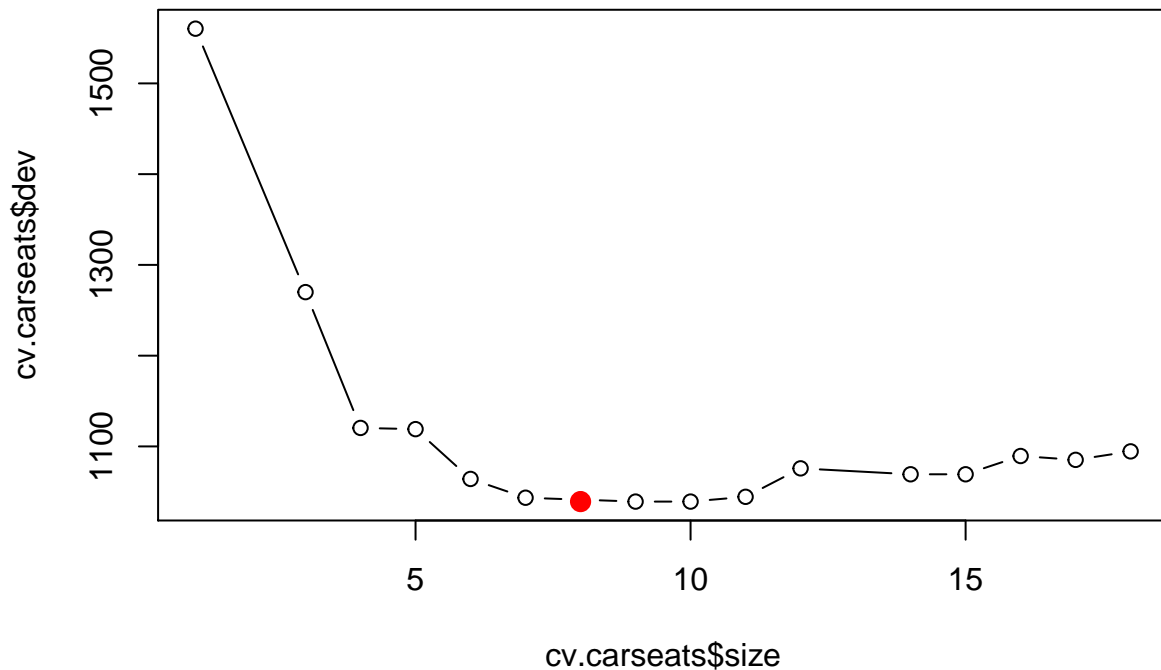
In this case, we obtain about 4.15 as test MSE.

c.

In this section, we use cross-validation to know which size is optimal for the tree complexity.

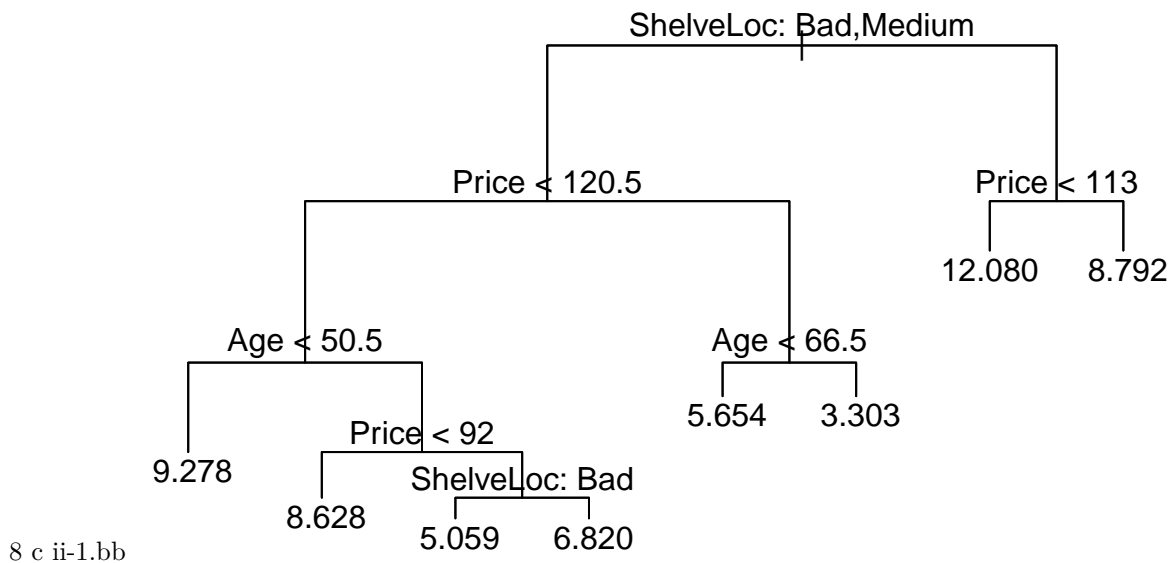
```
cv.carseats <- cv.tree(tree.carseats)
plot(cv.carseats$size, cv.carseats$dev, type = "b")
tree.min <- which.min(cv.carseats$dev)
points(tree.min, cv.carseats$dev[tree.min], col = "red", cex = 2, pch = 20)
```

8 c i-1.bb



We can observe that a size of 8 corresponds to the minimal deviance. Hence, we draw the tree with `size = 8` and we compute the test MSE for this tree.

```
prune.carseats = prune.tree(tree.carseats, best = 8)
plot(prune.carseats)
text(prune.carseats, pretty = 0)
```



```
pred = predict(prune.carseats, newdata = Carseats.test)
mean((pred - Carseats.test$Sales)^2)
```

```
## [1] 5.09085
```

We can see that pruning the tree is not helpful in this case since the test MSE increased after pruning from 4.15 to about 5.09. However, the tree seems way more human readable than the first one and presents the most important variables in the same order.

d.

```
library(randomForest)

## randomForest 4.6-14
## Type rfNews() to see new features/changes/bug fixes.
bag.carseats <- randomForest(Sales ~ ., data = Carseats.train, mtry = 10, ntree = 500, importance = TRUE)
pred.bag <- predict(bag.carseats, newdata = Carseats.test)
mean((pred.bag - Carseats.test$Sales)^2)

## [1] 2.633915
```

As we can see, using bagging decreased the test MSE to about 2.63 which means it outperforms the previous two models.

```
importance(bag.carseats)
```

	%IncMSE	IncNodePurity
CompPrice	16.9874366	126.852848
Income	3.8985402	78.314126
Advertising	16.5698586	123.702901
Population	0.6487058	62.328851
Price	55.3976775	514.654890
ShelveLoc	42.7849818	319.133777
Age	20.5135255	185.582077
Education	3.4615211	42.253410
Urban	-2.5125087	8.700009
US	7.3586645	18.180651

We can observe also that the most important variables are ShelveLoc, Price and Age which were also the most important variables for our previous models.

e.

```
mse.vec <- NA
for (a in 1:10){
  rf.carseats <- randomForest(Sales ~ ., data=Carseats.train,
                             mtry=a, ntree=500, importance=TRUE)
  rf.pred <- predict(rf.carseats, Carseats.test)
  mse.vec[a] <- mean((Carseats.test$Sales - rf.pred)^2)
}

# Number of variables used in the best model
which.min(mse.vec)
```

```
## [1] 10

# Test MSE corresponding to the best model
mse.vec[which.min(mse.vec)]
```

```
## [1] 2.56175
```

After applying random forest, we observe that the best model having to the lowest test MSE 2.56 corresponds to a tree using 10 predictors.

```
# Most important variables corresponding to the best model
rf.carseats <- randomForest(Sales ~ ., data = Carseats,
                             mtry=9, ntree=500, importance=TRUE)
importance(rf.carseats)
```

##		%IncMSE	IncNodePurity
##	CompPrice	37.4090331	328.33305
##	Income	11.6101066	167.93255
##	Advertising	29.6279922	235.62302
##	Population	-0.5069106	104.08331
##	Price	81.7594068	899.59222
##	ShelveLoc	86.8711457	972.32922
##	Age	27.2728885	286.99161
##	Education	2.8298995	84.28843
##	Urban	-1.8371710	12.60837
##	US	5.2177662	16.76991

Our best model shows also that the most important predictors in order are **ShelveLoc**, **Price**, **CompPrice**, **Advertising** and **Age**.

From the experiments above, we can conclude that the number of predictors **m** helps reduce the test MSE and hence, helps predicting our data with higher accuracy which explains why pruning the same tree in **b** generated a higher test MSE. However, there is a trade-off between the model's **accuracy** and **complexity** which makes the model more accurate but less readable and more complex.