

# DS502- HW3

*Mahdi Alouane and Rahul Pande*

## 1. (10 points) Section 6.8, page 259, question 2

- (a) iii. holds true. Lasso puts a budget constraint on the parameters which decreases the model variance and it reduces overfitting. However, when we put a constraint, the model bias increases. From the bias-variance trade-off concept we can say that the lasso regression will give better prediction when its increase in bias is less than its decrease in variance.
- (b) iii. holds true. Like Lasso above, Ridge also puts a budget constraint on the parameters which decreases the model variance and it reduces overfitting. However, when we put a constraint, the model bias increases. Similarly as above, from the bias-variance trade-off concept we can say that the Ridge regression will give better prediction when its increase in bias is less than its decrease in variance.
- (c) ii. holds true. Since non-linear methods are more flexible, they have higher variance than least squares regression but lower bias. Again, from the bias variance trade-off, if the increase on model variance is less than the decrease in bias, then non-linear model will have better prediction accuracy.

## 2. (20 points) Section 6.8, page 264, question 11

(a)

```
library(MASS)
library(leaps)
library(glmnet)
```

```
## Loading required package: Matrix
## Loading required package: foreach
## Loaded glmnet 2.0-16
```

```
attach(Boston)
```

```
colSums(sapply(Boston, is.na))
```

```
##      crim      zn      indus      chas      nox      rm      age      dis      rad
##         0         0         0         0         0         0         0         0         0
##      tax ptratio  black  lstat  medv
##         0         0         0         0         0
```

```
# No NAs in the dataset
```

```
# k-fold cross validation
```

```
k = 10
```

```
n = dim(Boston)[1]
```

```
p = dim(Boston)[2]-1
```

```
set.seed(123)
```

```
folds = sample(rep(1:k, length = nrow(Boston)), replace = T)
```

```

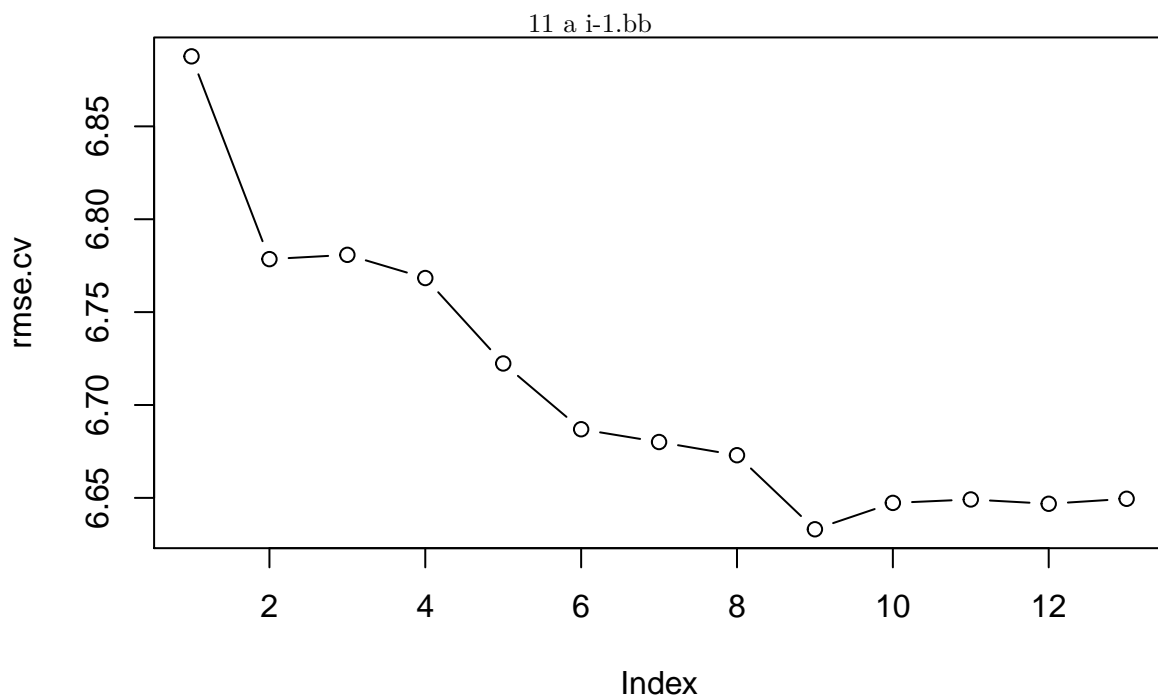
form.subset = as.formula("crim ~ .")

cv.errors = matrix(NA, k, p)

for (i in 1:k) {
  subset.fit <- regsubsets(form.subset, Boston[folds!=i, ], nvmax = p)
  for (n.subset in 1:p) {
    m.mat <- model.matrix(form.subset, Boston[folds==i, ])
    best.coef <- coef(subset.fit, id=n.subset)
    pred <- m.mat[, names(best.coef)] %*% best.coef
    # mean squared error
    error = mean((Boston[folds==i, ]$crim - pred)^2)
    cv.errors[i,n.subset] = error
  }
}

# root mean squared values
rmse.cv = sqrt(apply(cv.errors, 2, mean))
plot(rmse.cv, type = "b")

```



```

which(rmse.cv == min(rmse.cv))

## [1] 9

rmse.cv[which(rmse.cv == min(rmse.cv))]

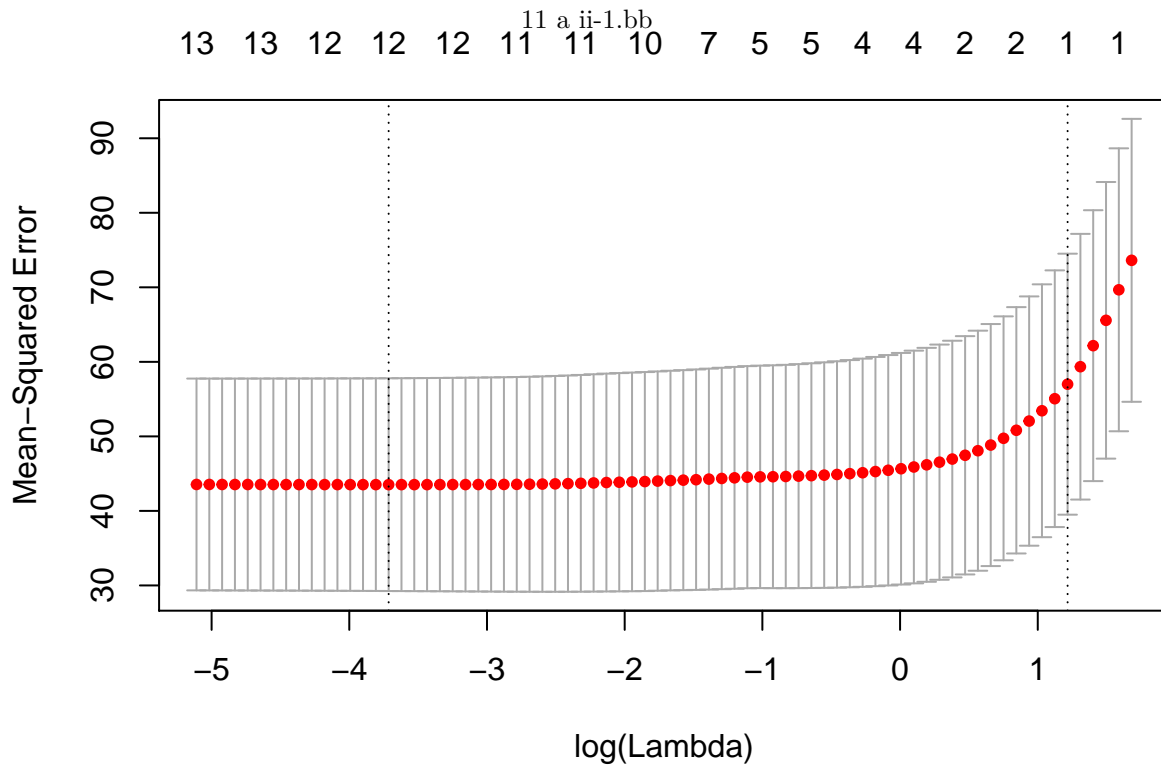
## [1] 6.633116

# Lasso regression

X = model.matrix(crim ~ ., data = Boston)[, -1]
y = Boston$crim
cv.lasso = cv.glmnet(X, y, type.measure = "mse", nfolds = 10)

```

```
plot(cv.lasso)
```



```
coef(cv.lasso)[,1]
```

```
## (Intercept)      zn      indus      chas      nox      rm
##  1.4186415  0.0000000  0.0000000  0.0000000  0.0000000  0.0000000
##      age      dis      rad      tax      ptratio  black
##  0.0000000  0.0000000  0.2298449  0.0000000  0.0000000  0.0000000
##      lstat      medv
##  0.0000000  0.0000000
```

```
# One standard error lamda to avoid overfitting
```

```
chosen.lambda = cv.lasso$lambda.1se
```

```
# root mean square error for the chosen lamdba
```

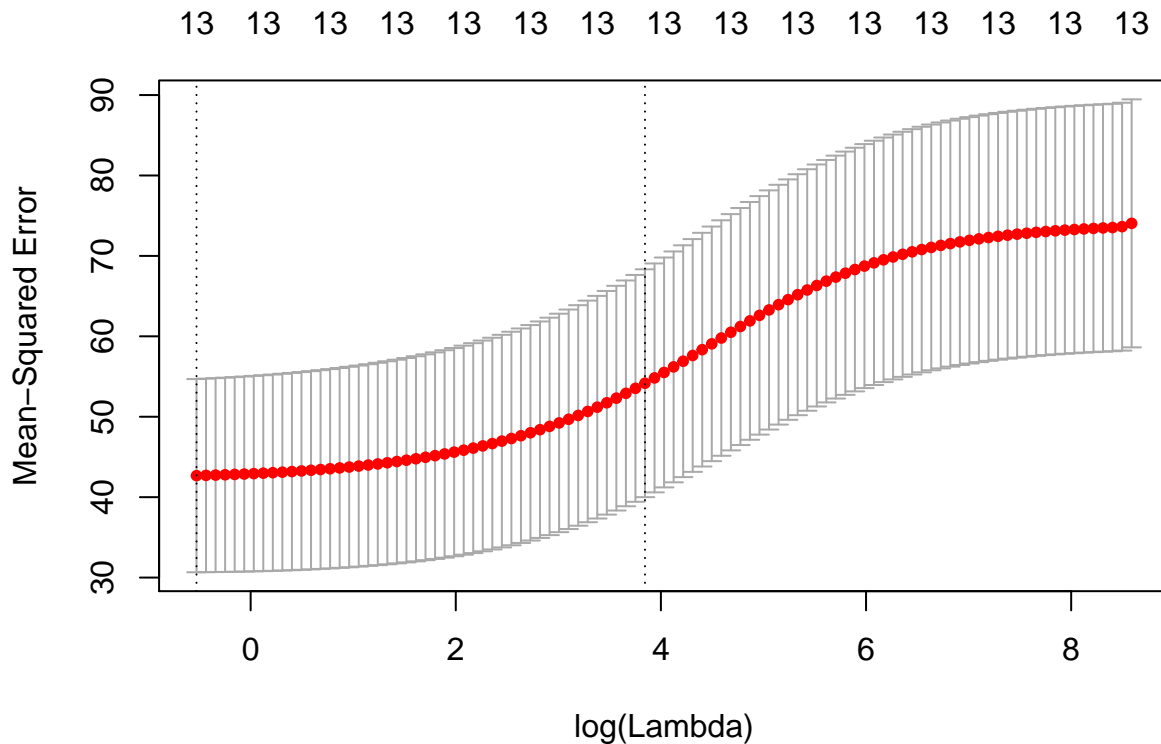
```
sqrt(cv.lasso$cvm[cv.lasso$lambda == chosen.lambda])
```

```
## [1] 7.549995
```

```
cv.ridge = cv.glmnet(X, y, type.measure = "mse", alpha = 0, nfolds = 10)
```

```
plot(cv.ridge)
```

11 a iii-1.bb



```
coef(cv.ridge)
```

```
## 14 x 1 sparse Matrix of class "dgCMatrix"
##              1
## (Intercept)  0.757566248
## zn          -0.002519179
## indus        0.036530724
## chas        -0.259774554
## nox          2.423435801
## rm          -0.169106780
## age          0.007845252
## dis         -0.125063596
## rad          0.067739853
## tax          0.002972047
## ptratio      0.093676787
## black       -0.003748456
## lstat        0.049092196
## medv        -0.032058198
```

```
# One standard error lamda to avoid overfitting
```

```
chosen.lambda = cv.ridge$lambda.1se
```

```
# root mean square error for the chosen lamdba
```

```
sqrt(cv.ridge$cvm[cv.ridge$lambda == chosen.lambda])
```

```
## [1] 7.359946
```

```
library(pls)
```

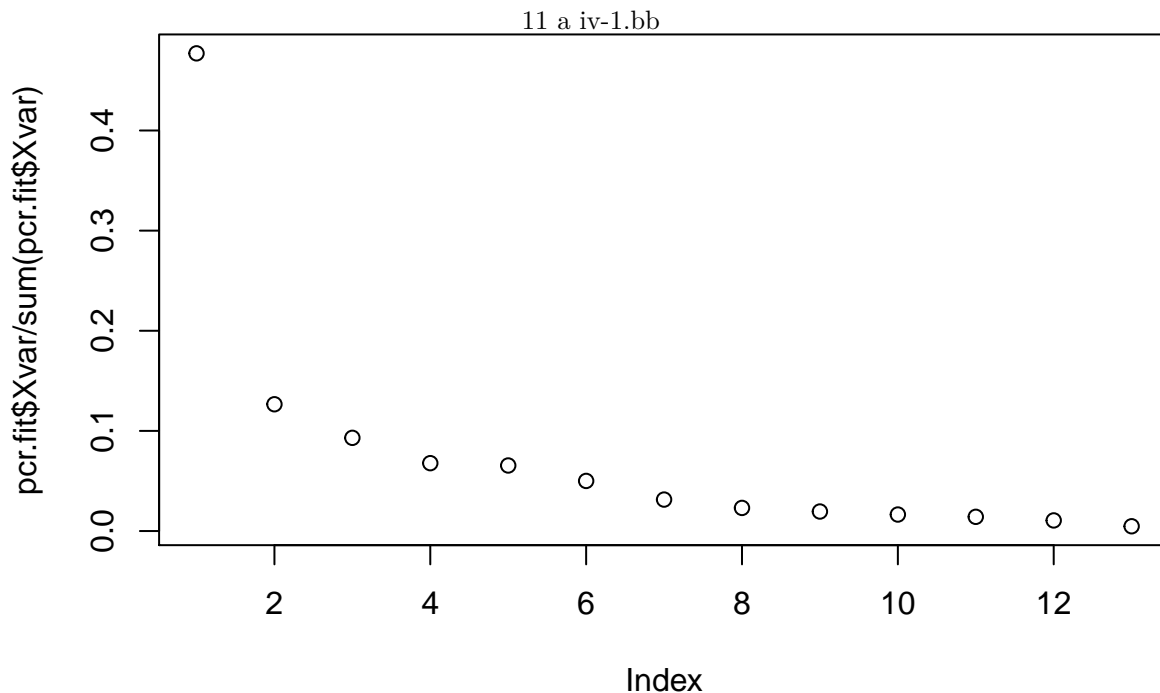
```
##
```

```
## Attaching package: 'pls'
```


```
## The following object is masked from 'package:stats':
##
##   loadings
pcr.fit = pcr(crim ~ ., data = Boston, scale = TRUE, validation = "CV", segments= 10)
summary(pcr.fit)

## Data:      X dimension: 506 13
## Y dimension: 506 1
## Fit method: svdpc
## Number of components considered: 13
##
## VALIDATION: RMSEP
## Cross-validated using 10 random segments.
##      (Intercept)  1 comps  2 comps  3 comps  4 comps  5 comps  6 comps
## CV              8.61    7.189    7.185    6.762    6.770    6.769    6.784
## adjCV           8.61    7.187    7.183    6.757    6.764    6.764    6.778
##      7 comps  8 comps  9 comps 10 comps 11 comps 12 comps 13 comps
## CV          6.789    6.653    6.696    6.680    6.673    6.626    6.563
## adjCV       6.782    6.645    6.686    6.671    6.663    6.615    6.551
##
## TRAINING: % variance explained
##      1 comps  2 comps  3 comps  4 comps  5 comps  6 comps  7 comps
## X          47.70    60.36    69.67    76.45    82.99    88.00    91.14
## crim       30.69    30.87    39.27    39.61    39.61    39.86    40.14
##      8 comps  9 comps 10 comps 11 comps 12 comps 13 comps
## X          93.45    95.40    97.04    98.46    99.52    100.0
## crim       42.47    42.55    42.78    43.04    44.13    45.4

# variance explanation
plot(pcr.fit$Xvar/sum(pcr.fit$Xvar))
```



```
# validation plot
validationplot(pcr.fit, val.type='MSEP')
```



```
ncomp = 4

set.seed(123)
folds = sample(rep(1:k, length = nrow(Boston)), replace = T)

cv.errors = rep(0, k)

for (i in 1:k) {
  pcr.fit <- pcr(crim ~ ., data = Boston[folds != i, ], scale = TRUE)
  pred = predict(pcr.fit, Boston[folds == i, ], ncomp=ncomp)
  error = mean((Boston[folds == i, ]$crim - pred)^2)
  cv.errors[i] = error
}
sqrt(mean(cv.errors))
```

```
## [1] 6.875975
```

(b)

```
results <- rbind(
  c("Best Subset", 6.633116, 9),
  c("Lasso Regression", 7.549995, 1),
  c("Ridge Regression", 7.359946, 13),
  c("PCR", 6.875975, 4)
)
colnames(results) <- c("Method", "MSE", "# predictors")
```

```
knitr::kable(results)
```

Method	MSE	# predictors
Best Subset	6.633116	9
Lasso Regression	7.549995	1
Ridge Regression	7.359946	13
PCR	6.875975	4

From the above table, we chose PCR model with 4 predictors as it has the mse very close to the lowest mse and is simpler model than the best subset model, since it has 4 predictors against the 9 predictors and thus has lower chances of overfitting the data. Second choice would be Best Subset model with 9 predictors as it has the lowest cross validation mse and has less number of predictors than Ridge Regression. Lasso Regression here seems to be underfit in this case.

- (c) No. PCR has only 4 predictors since from the graph, after 4 predictors, adding another predictor does not increase the explained variance a lot. Hence we have taken only 4 predictors. We can also have criterion like taking  $n$  components which explain at least  $x\%$  of the variance.

### 3. (10 points) Section 7.9, Page 298, question 3

We have  $b_1(X) = X, b_2(X) = (X - 1)^2 * I(X \geq 1)$

For  $X \geq 1, I(X \geq 1) = 1$ , and  $X < 1, I(X \geq 1) = 0$

Substituting  $\hat{\beta}_0 = 1, \hat{\beta}_1 = 1, \hat{\beta}_2 = 2$  in

$$Y = \beta_0 + \beta_1 b_1(X) + \beta_2 b_2(X) + \epsilon$$

We get,

$$\hat{Y} = 1 + b_1(X) - 2b_2(X)$$

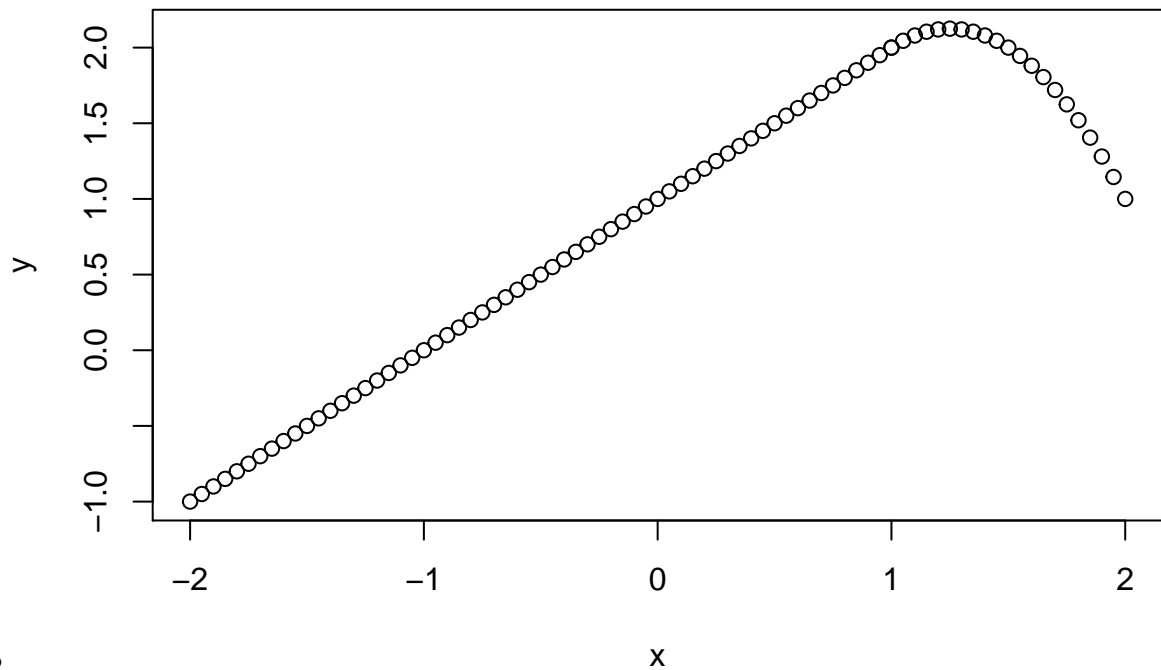
For  $X \geq 1, \hat{Y} = 1 + X - 2(X - 1)^2 = -1 + 5X - 2X^2$ , and for  $X < 1, \hat{Y} = 1 + X$

```
x_lower = seq(-2, 1, by = 0.05)
x_upper = seq(1, 2, by = 0.05)

y_lower = 1 + x_lower
y_upper = -1 + 5 * x_upper - 2 * (x_upper ^ 2)

x <- c(x_lower, x_upper)
y <- c(y_lower, y_upper)

plot(x,y)
```



2-1.bb

$Y = 1 + X$ , for  $X < 1$  Now, for  $X = 0$ ,  $Y = 1$

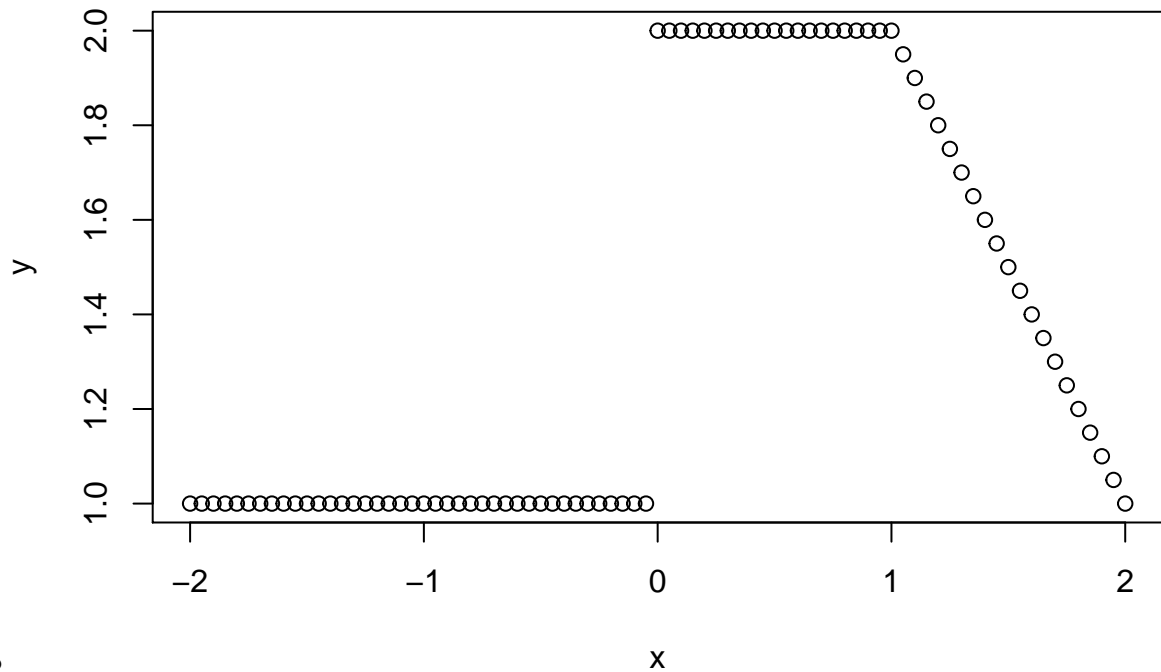
Therefore y-intercept is 1. Slope is 1 when  $X < 1$  and by taking derivative for  $Y$  where  $X \geq 1$ , we get slope as  $5 - 4X$

#### 4. (10 points) Section 7.9, Page 298, question 4

Similar to from above, we can split the function into multiple domains. Since there are a lot of cuts in this, we use the I function in R to enforce the conditions on x. It is as below.

```
x = seq(-2, 2, 0.05)
y = 1 + 1 * I(x <= 2 & x >= 0) - (x-1) * I(x <= 2 & x >= 1) + 3 * (x-3) * I(x <= 4 & x >= 3) + I(x <=
plot(x, y)
```





4-1.bb

```
# y-intercept
y[which(x==0)]
```

```
## [1] 2
```

The y-intercept is 2 (`y[which(x==0)]`). Slope is -1 for  $1 \leq X \leq 2$  and 0 for  $-2 \leq X < 0$  and  $0 < X \leq 1$ . The function is discontinuous at  $x=0$

5. (10 points) Section 7.9, Page 299, question 6
6. (20 points) Section 7.9, Page 299, question 7
7. (10 points) Section 8.4, Page 332, question 1
8. (10 points) Section 8.4, Page 333-334, question 8